

04_supervised_learning

September 4, 2025

1 04 - Supervised Learning (Classification)

This notebook includes: - Target variable definition pass/fail - Feature engineering and encoding - Multiple classification algorithms - Hyperparameter tuning - Cross-validation - Model interpretation

```
[2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import cross_val_score
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import RandomizedSearchCV

try:
    from xgboost import XGBClassifier

    print("XGBoost successfully imported!")
except ImportError:
    print("XGBoost not installed. Install with: pip install xgboost")
    XGBClassifier = None
import warnings

warnings.filterwarnings("ignore")
```

XGBoost successfully imported!

```
[3]: df = pd.read_pickle(r"../data/processed/cleaned_dataset.pkl")
df.drop(columns=["grade_average"], inplace=True)
```

```

[4]: # Handle different data types properly after preprocessing

# Apply encoding
df_encoded = df.copy()

# 1. ORDINAL VARIABLES - Convert ordered categoricals to numeric (preserve
↳ order)
ordinal_vars = [
    "Medu",
    "Fedu",
    "traveltime",
    "studytime",
    "famrel",
    "freetime",
    "goout",
    "Dalc",
    "Walc",
    "health",
]
for var in ordinal_vars:
    if var in df_encoded.columns:
        df_encoded[var] = df_encoded[
            var
        ].cat.codes # Converts to 0,1,2,3,4 maintaining order
        print(f"Ordinal encoded {var}")

# 2. BINARY VARIABLES - Convert booleans to 0/1
binary_vars = [
    "schoolsup",
    "famsup",
    "paid",
    "activities",
    "nursery",
    "higher",
    "internet",
    "romantic",
]
for var in binary_vars:
    if var in df_encoded.columns:
        df_encoded[var] = df_encoded[var].astype(int) # True->1, False->0
        print(f"Boolean to binary: {var}")

# 3. NOMINAL CATEGORICALS - Binary mapping for 2-level, one-hot for multi-level
# Binary mappings for 2-level categoricals
binary_mappings = {
    "school": {"GP": 1, "MS": 0},
    "sex": {"F": 1, "M": 0},

```

```

    "address": {"U": 1, "R": 0},
    "famsize": {"GT3": 1, "LE3": 0},
    "Pstatus": {"T": 1, "A": 0},
}

# Apply binary encoding
for col, mapping in binary_mappings.items():
    if col in df_encoded.columns:
        df_encoded[col] = df_encoded[col].map(mapping)
        print(f"Binary encoded {col}")

# Multi-level categorical columns for one-hot encoding
categorical_cols = ["Mjob", "Fjob", "reason", "guardian"]

# One-hot encode categorical columns
df_encoded = pd.get_dummies(df_encoded, columns=categorical_cols,
                             drop_first=True)
print(f"One-hot encoded: {categorical_cols}")
print(f"New shape after encoding: {df_encoded.shape}")

# 4. NUMERIC VARIABLES already handled (age, failures, absences, G1, G2, G3)

# Remove target leakage
features_to_drop = ["G3"]
X = df_encoded.drop(columns=features_to_drop + ["pass_fail"])
y = df_encoded["pass_fail"]

# Train/test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

# Scale features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Results
print(f"\nFinal results:")
print(f"Training shape: {X_train_scaled.shape}")
print(f"Test shape: {X_test_scaled.shape}")
print(f"Pass rate - Train: {y_train.mean():.1%}, Test: {y_test.mean():.1%}")
print(f"All features are now numeric and scaled.")

```

Ordinal encoded Medu
Ordinal encoded Fedu
Ordinal encoded traveltime
Ordinal encoded studytime

Ordinal encoded famrel
Ordinal encoded freetime
Ordinal encoded goout
Ordinal encoded Dalc
Ordinal encoded Walc
Ordinal encoded health
Boolean to binary: schoolsup
Boolean to binary: famsup
Boolean to binary: paid
Boolean to binary: activities
Boolean to binary: nursery
Boolean to binary: higher
Boolean to binary: internet
Boolean to binary: romantic
Binary encoded school
Binary encoded sex
Binary encoded address
Binary encoded famsize
Binary encoded Pstatus
One-hot encoded: ['Mjob', 'Fjob', 'reason', 'guardian']
New shape after encoding: (649, 44)

Final results:

Training shape: (519, 42)
Test shape: (130, 42)
Pass rate - Train: 84.6%, Test: 84.6%
All features are now numeric and scaled.

Final results:

Training shape: (519, 42)
Test shape: (130, 42)
Pass rate - Train: 84.6%, Test: 84.6%
All features are now numeric and scaled.

```
[5]: # Create two feature sets from the scaled data
print("Available features:", X.columns.tolist())

# Set 1: Keep all current features (includes G1, G2)
X_with_grades = X.copy()
X_train_with_grades = X_train_scaled.copy()
X_test_with_grades = X_test_scaled.copy()

# Set 2: Remove G1 and G2 from the scaled data
grade_columns = ["G1", "G2"]
grade_indices = [X.columns.get_loc(col) for col in grade_columns if col in X.
                  columns]
```

```

X_without_grades = X.drop(columns=grade_columns)
X_train_without_grades = np.delete(X_train_scaled, grade_indices, axis=1)
X_test_without_grades = np.delete(X_test_scaled, grade_indices, axis=1)

# Results
print(f"\nFeature sets created:")
print(
    f"WITH G1, G2 - Train: {X_train_with_grades.shape}, Test:␣
↪{X_test_with_grades.shape}"
)
print(
    f"WITHOUT G1, G2 - Train: {X_train_without_grades.shape}, Test:␣
↪{X_test_without_grades.shape}"
)
print(f"Features removed: {grade_columns}")

```

Available features: ['school', 'sex', 'age', 'address', 'famsize', 'Pstatus', 'Medu', 'Fedu', 'traveltime', 'studytime', 'failures', 'schoolsup', 'famsup', 'paid', 'activities', 'nursery', 'higher', 'internet', 'romantic', 'famrel', 'freetime', 'goout', 'Dalc', 'Walc', 'health', 'absences', 'G1', 'G2', 'attendance_proxy', 'Mjob_health', 'Mjob_other', 'Mjob_services', 'Mjob_teacher', 'Fjob_health', 'Fjob_other', 'Fjob_services', 'Fjob_teacher', 'reason_home', 'reason_other', 'reason_reputation', 'guardian_mother', 'guardian_other']

Feature sets created:

WITH G1, G2 - Train: (519, 42), Test: (130, 42)

WITHOUT G1, G2 - Train: (519, 40), Test: (130, 40)

Features removed: ['G1', 'G2']

2 Function-Based Model Implementation

2.1 Logistic Regression

```

[6]: def logistic_regression_base(X_train, X_test, title):
    """Base Logistic Regression Model"""
    print(f"LOGISTIC REGRESSION MODEL {title.upper()}")

    # Initialize and train the model
    log_reg = LogisticRegression(random_state=42)
    log_reg.fit(X_train, y_train)

    # Make predictions on TEST data
    y_pred_lr = log_reg.predict(X_test)

    # Evaluate the model
    print("\nClassification Report:")

```

```

print(classification_report(y_test, y_pred_lr))

# Logistic Regression metrics
lr_accuracy = accuracy_score(y_test, y_pred_lr)
lr_precision = precision_score(y_test, y_pred_lr)
lr_recall = recall_score(y_test, y_pred_lr)
lr_f1 = f1_score(y_test, y_pred_lr)

# Confusion Matrix
print("Confusion Matrix:")
cm = confusion_matrix(y_test, y_pred_lr)
plt.figure(figsize=(8, 6))
lr_color = "#0d47a1"
sns.heatmap(
    cm,
    annot=True,
    fmt="d",
    cmap=sns.light_palette(lr_color, as_cmap=True),
)
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title(f"Logistic Regression - {title}")
plt.show()

return lr_accuracy, lr_precision, lr_recall, lr_f1

def logistic_regression_tuned(X_train, X_test, title):
    """Hyperparameter Tuned Logistic Regression Model"""
    print(f"HYPERPARAMETER TUNED LOGISTIC REGRESSION MODEL {title.upper()}")

    # Define the parameter grid
    param_grid = {
        "C": [0.001, 0.01, 0.1, 1, 10, 100],
        "penalty": ["l1", "l2"],
        "solver": ["liblinear", "saga"],
    }

    # Initialize GridSearchCV
    grid_search = GridSearchCV(
        LogisticRegression(random_state=42, max_iter=1000),
        param_grid,
        cv=5,
        scoring="f1",
        n_jobs=-1,
    )

```

```

# Fit the grid search to the data
grid_search.fit(X_train, y_train)

# Get the best model
best_log_reg = grid_search.best_estimator_

# Print the best parameters
print("Best Hyperparameters for Logistic Regression:")
print(grid_search.best_params_)

# Make predictions with the improved model
y_pred_lr = best_log_reg.predict(X_test)

# Evaluate the improved model
print("\nClassification Report for Improved Logistic Regression:")
print(classification_report(y_test, y_pred_lr))

# Tuned Logistic Regression metrics
lr_tuned_accuracy = accuracy_score(y_test, y_pred_lr)
lr_tuned_precision = precision_score(y_test, y_pred_lr)
lr_tuned_recall = recall_score(y_test, y_pred_lr)
lr_tuned_f1 = f1_score(y_test, y_pred_lr)

# Confusion Matrix for the improved model
print("\nConfusion Matrix for Improved Model:")
cm_improved = confusion_matrix(y_test, y_pred_lr)
plt.figure(figsize=(8, 6))
lr_tuned_color = "#42a5f5"
sns.heatmap(
    cm_improved,
    annot=True,
    fmt="d",
    cmap=sns.light_palette(lr_tuned_color, as_cmap=True),
)
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title(f"Tuned Logistic Regression - {title}")
plt.show()

return lr_tuned_accuracy, lr_tuned_precision, lr_tuned_recall, lr_tuned_f1

```

2.2 Random Forest

```

[7]: def random_forest_base(X_train, X_test, title):
    """Base Random Forest Model"""
    print(f"RANDOM FOREST MODEL {title.upper()}")

```

```

# Initialize and train the model
rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)

# Make predictions
y_pred_rf = rf.predict(X_test)

# Evaluate the model
print("\nClassification Report:")
print(classification_report(y_test, y_pred_rf))

# Random Forest metrics
rf_accuracy = accuracy_score(y_test, y_pred_rf)
rf_precision = precision_score(y_test, y_pred_rf)
rf_recall = recall_score(y_test, y_pred_rf)
rf_f1 = f1_score(y_test, y_pred_rf)

# Confusion Matrix
print("Confusion Matrix:")
cm_rf = confusion_matrix(y_test, y_pred_rf)
plt.figure(figsize=(8, 6))
rf_color = "#1b5e20"
sns.heatmap(
    cm_rf,
    annot=True,
    fmt="d",
    cmap=sns.light_palette(rf_color, as_cmap=True),
)
plt.title(f"Random Forest - {title}")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

return rf_accuracy, rf_precision, rf_recall, rf_f1

def random_forest_tuned(X_train, X_test, title):
    """Hyperparameter Tuned Random Forest Model"""
    print(f"HYPERPARAMETER TUNED RANDOM FOREST MODEL {title.upper()}")

    # Define the parameter distribution for Randomized Search
    param_dist = {
        "n_estimators": [100, 200, 300],
        "max_depth": [10, 20, 30, None],
        "min_samples_split": [2, 5, 10],
        "min_samples_leaf": [1, 2, 4],
        "bootstrap": [True, False],
    }

```



```

}

# Initialize RandomizedSearchCV
random_search = RandomizedSearchCV(
    RandomForestClassifier(random_state=42),
    param_distributions=param_dist,
    n_iter=20, # Try 20 different combinations
    cv=5,
    scoring="f1",
    n_jobs=-1,
    random_state=42,
)

random_search.fit(X_train, y_train)

# Get the best model
best_rf = random_search.best_estimator_

print("Best Hyperparameters for Random Forest:")
print(random_search.best_params_)

# Make predictions with the tuned model
y_pred_rf = best_rf.predict(X_test)

# Evaluate the tuned model
print("\nClassification Report for Tuned Random Forest:")
print(classification_report(y_test, y_pred_rf))

# Tuned Random Forest metrics
rf_tuned_accuracy = accuracy_score(y_test, y_pred_rf)
rf_tuned_precision = precision_score(y_test, y_pred_rf)
rf_tuned_recall = recall_score(y_test, y_pred_rf)
rf_tuned_f1 = f1_score(y_test, y_pred_rf)

# Confusion Matrix for the tuned model
print("\nConfusion Matrix for Tuned Model:")
cm_rf_tuned = confusion_matrix(y_test, y_pred_rf)
plt.figure(figsize=(8, 6))
rf_tuned_color = "#66bb6a"
sns.heatmap(
    cm_rf_tuned,
    annot=True,
    fmt="d",
    cmap=sns.light_palette(rf_tuned_color, as_cmap=True),
)
plt.xlabel("Predicted")
plt.ylabel("Actual")

```

```
plt.title(f"Tuned Random Forest - {title}")
plt.show()

return rf_tuned_accuracy, rf_tuned_precision, rf_tuned_recall, rf_tuned_f1
```

2.3 Gradient Boosting Machine

```
[8]: def gradient_boosting_base(X_train, X_test, title):
    """Base Gradient Boosting Model"""
    print(f"GRADIENT BOOSTING MODEL {title.upper()}")

    gbm = GradientBoostingClassifier(random_state=42)
    gbm.fit(X_train, y_train)

    y_pred_gbm = gbm.predict(X_test)

    # Evaluate the model
    print("\nClassification Report:")
    print(classification_report(y_test, y_pred_gbm))

    # Gradient Boosting metrics
    gbm_accuracy = accuracy_score(y_test, y_pred_gbm)
    gbm_precision = precision_score(y_test, y_pred_gbm)
    gbm_recall = recall_score(y_test, y_pred_gbm)
    gbm_f1 = f1_score(y_test, y_pred_gbm)

    # Confusion Matrix
    print("Confusion Matrix:")
    cm_gbm = confusion_matrix(y_test, y_pred_gbm)
    plt.figure(figsize=(8, 6))
    gbm_color = "#4a148c"
    sns.heatmap(
        cm_gbm,
        annot=True,
        fmt="d",
        cmap=sns.light_palette(gbm_color, as_cmap=True),
    )
    plt.title(f"Gradient Boosting - {title}")
    plt.xlabel("Predicted")
    plt.ylabel("Actual")
    plt.show()

    return gbm_accuracy, gbm_precision, gbm_recall, gbm_f1

def gradient_boosting_tuned(X_train, X_test, title):
    """Hyperparameter Tuned Gradient Boosting Model"""
```

```

print(f"TUNED GRADIENT BOOSTING MODEL {title.upper()}")

# Define the parameter distribution for Randomized Search for GBM
param_dist_gbm = {
    "n_estimators": [100, 200, 300],
    "learning_rate": [0.01, 0.1, 0.2],
    "max_depth": [3, 5, 7],
    "min_samples_split": [2, 5, 10],
    "min_samples_leaf": [1, 2, 4],
    "subsample": [0.8, 0.9, 1.0],
}

# Initialize RandomizedSearchCV for GBM
random_search_gbm = RandomizedSearchCV(
    GradientBoostingClassifier(random_state=42),
    param_distributions=param_dist_gbm,
    n_iter=20, # Try 20 different combinations
    cv=5,
    scoring="f1",
    n_jobs=-1,
    random_state=42,
)

# Fit the random search to the data
random_search_gbm.fit(X_train, y_train)

# Get the best model
best_gbm = random_search_gbm.best_estimator_

# Print the best parameters
print("Best Hyperparameters for Gradient Boosting:")
print(random_search_gbm.best_params_)

# Make predictions with the tuned model
y_pred_gbm = best_gbm.predict(X_test)

# Evaluate the tuned model
print("\nClassification Report for Tuned Gradient Boosting:")
print(classification_report(y_test, y_pred_gbm))

# Tuned Gradient Boosting metrics
gbm_tuned_accuracy = accuracy_score(y_test, y_pred_gbm)
gbm_tuned_precision = precision_score(y_test, y_pred_gbm)
gbm_tuned_recall = recall_score(y_test, y_pred_gbm)
gbm_tuned_f1 = f1_score(y_test, y_pred_gbm)

# Confusion Matrix for the tuned model

```

```

print("\nConfusion Matrix for Tuned GBM:")
cm_gbm_tuned = confusion_matrix(y_test, y_pred_gbm)
plt.figure(figsize=(8, 6))
gbm_tuned_color = "#9575cd"
sns.heatmap(
    cm_gbm_tuned,
    annot=True,
    fmt="d",
    cmap=sns.light_palette(gbm_tuned_color, as_cmap=True),
)
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title(f"Tuned Gradient Boosting - {title}")
plt.show()

return gbm_tuned_accuracy, gbm_tuned_precision, gbm_tuned_recall,
↪gbm_tuned_f1

```

2.4 Support Vector Machine (SVM)

```

[9]: def svm_base(X_train, X_test, title):
    """Base SVM Model"""
    from sklearn.svm import SVC

    print(f"SVM MODEL {title.upper()}")

    # Initialize and train the model
    svm = SVC(
        random_state=42, probability=True
    ) # probability=True for better integration
    svm.fit(X_train, y_train)

    # Make predictions
    y_pred_svm = svm.predict(X_test)

    # Evaluate the model
    print("\nClassification Report:")
    print(classification_report(y_test, y_pred_svm))

    # SVM metrics
    svm_accuracy = accuracy_score(y_test, y_pred_svm)
    svm_precision = precision_score(y_test, y_pred_svm)
    svm_recall = recall_score(y_test, y_pred_svm)
    svm_f1 = f1_score(y_test, y_pred_svm)

    # Confusion Matrix
    print("Confusion Matrix:")

```

```

cm_svm = confusion_matrix(y_test, y_pred_svm)
plt.figure(figsize=(8, 6))
svm_color = "#d32f2f" # Red color
sns.heatmap(
    cm_svm,
    annot=True,
    fmt="d",
    cmap=sns.light_palette(svm_color, as_cmap=True),
)
plt.title(f"SVM - {title}")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

return svm_accuracy, svm_precision, svm_recall, svm_f1

def svm_tuned(X_train, X_test, title):
    """Hyperparameter Tuned SVM Model"""
    from sklearn.svm import SVC

    print(f"TUNED SVM MODEL {title.upper()}")

    # Define the parameter distribution for Randomized Search
    param_dist_svm = {
        "C": [0.1, 1, 10, 100, 1000],
        "gamma": ["scale", "auto", 0.001, 0.01, 0.1, 1],
        "kernel": ["rbf", "poly", "sigmoid"],
        "degree": [2, 3, 4], # Only for poly kernel
    }

    # Initialize RandomizedSearchCV for SVM
    random_search_svm = RandomizedSearchCV(
        SVC(random_state=42, probability=True),
        param_distributions=param_dist_svm,
        n_iter=20, # Try 20 different combinations
        cv=5,
        scoring="f1",
        n_jobs=-1,
        random_state=42,
    )

    # Fit the random search to the data
    random_search_svm.fit(X_train, y_train)

    # Get the best model
    best_svm = random_search_svm.best_estimator_

```

```

# Print the best parameters
print("Best Hyperparameters for SVM:")
print(random_search_svm.best_params_)

# Make predictions with the tuned model
y_pred_svm = best_svm.predict(X_test)

# Evaluate the tuned model
print("\nClassification Report for Tuned SVM:")
print(classification_report(y_test, y_pred_svm))

# Tuned SVM metrics
svm_tuned_accuracy = accuracy_score(y_test, y_pred_svm)
svm_tuned_precision = precision_score(y_test, y_pred_svm)
svm_tuned_recall = recall_score(y_test, y_pred_svm)
svm_tuned_f1 = f1_score(y_test, y_pred_svm)

# Confusion Matrix for the tuned model
print("\nConfusion Matrix for Tuned SVM:")
cm_svm_tuned = confusion_matrix(y_test, y_pred_svm)
plt.figure(figsize=(8, 6))
svm_tuned_color = "#f44336" # Light red
sns.heatmap(
    cm_svm_tuned,
    annot=True,
    fmt="d",
    cmap=sns.light_palette(svm_tuned_color, as_cmap=True),
)
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title(f"Tuned SVM - {title}")
plt.show()

return svm_tuned_accuracy, svm_tuned_precision, svm_tuned_recall,
↪svm_tuned_f1

```

2.5 XGBoost (eXtreme Gradient Boosting)

```

[10]: def xgboost_base(X_train, X_test, title):
        """Base XGBoost Model"""
        if XGBClassifier is None:
            print("XGBoost not available. Please install with: pip install xgboost")
            return 0, 0, 0, 0

        print(f"XGBOOST MODEL {title.upper()}")

```

```

# Initialize and train the model
xgb = XGBClassifier(
    random_state=42,
    eval_metric="logloss", # Suppress warning
    verbosity=0, # Reduce output
)
xgb.fit(X_train, y_train)

# Make predictions
y_pred_xgb = xgb.predict(X_test)

# Evaluate the model
print("\nClassification Report:")
print(classification_report(y_test, y_pred_xgb))

# XGBoost metrics
xgb_accuracy = accuracy_score(y_test, y_pred_xgb)
xgb_precision = precision_score(y_test, y_pred_xgb)
xgb_recall = recall_score(y_test, y_pred_xgb)
xgb_f1 = f1_score(y_test, y_pred_xgb)

# Confusion Matrix
print("Confusion Matrix:")
cm_xgb = confusion_matrix(y_test, y_pred_xgb)
plt.figure(figsize=(8, 6))
xgb_color = "#e65100" # Orange color
sns.heatmap(
    cm_xgb,
    annot=True,
    fmt="d",
    cmap=sns.light_palette(xgb_color, as_cmap=True),
)
plt.title(f"XGBoost - {title}")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

return xgb_accuracy, xgb_precision, xgb_recall, xgb_f1

def xgboost_tuned(X_train, X_test, title):
    """Hyperparameter Tuned XGBoost Model"""
    if XGBClassifier is None:
        print("XGBoost not available. Please install with: pip install xgboost")
        return 0, 0, 0, 0

    print(f"TUNED XGBOOST MODEL {title.upper()}")

```

```

# Define the parameter distribution for Randomized Search
param_dist_xgb = {
    "n_estimators": [100, 200, 300],
    "learning_rate": [0.01, 0.1, 0.2],
    "max_depth": [3, 5, 7],
    "min_child_weight": [1, 3, 5],
    "subsample": [0.8, 0.9, 1.0],
    "colsample_bytree": [0.8, 0.9, 1.0],
    "reg_alpha": [0, 0.1, 0.5], # L1 regularization
    "reg_lambda": [1, 1.5, 2], # L2 regularization
}

# Initialize RandomizedSearchCV for XGBoost
random_search_xgb = RandomizedSearchCV(
    XGBClassifier(random_state=42, eval_metric="logloss", verbosity=0),
    param_distributions=param_dist_xgb,
    n_iter=20, # Try 20 different combinations
    cv=5,
    scoring="f1",
    n_jobs=-1,
    random_state=42,
)

# Fit the random search to the data
random_search_xgb.fit(X_train, y_train)

# Get the best model
best_xgb = random_search_xgb.best_estimator_

# Print the best parameters
print("Best Hyperparameters for XGBoost:")
print(random_search_xgb.best_params_)

# Make predictions with the tuned model
y_pred_xgb = best_xgb.predict(X_test)

# Evaluate the tuned model
print("\nClassification Report for Tuned XGBoost:")
print(classification_report(y_test, y_pred_xgb))

# Tuned XGBoost metrics
xgb_tuned_accuracy = accuracy_score(y_test, y_pred_xgb)
xgb_tuned_precision = precision_score(y_test, y_pred_xgb)
xgb_tuned_recall = recall_score(y_test, y_pred_xgb)
xgb_tuned_f1 = f1_score(y_test, y_pred_xgb)

```



```

# Confusion Matrix for the tuned model
print("\nConfusion Matrix for Tuned XGBoost:")
cm_xgb_tuned = confusion_matrix(y_test, y_pred_xgb)
plt.figure(figsize=(8, 6))
xgb_tuned_color = "#ff9800" # Light orange
sns.heatmap(
    cm_xgb_tuned,
    annot=True,
    fmt="d",
    cmap=sns.light_palette(xgb_tuned_color, as_cmap=True),
)
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title(f"Tuned XGBoost - {title}")
plt.show()

return xgb_tuned_accuracy, xgb_tuned_precision, xgb_tuned_recall,
↪xgb_tuned_f1

```

2.6 Model Execution and Comparison

```

[11]: # Initialize results dictionary for comparison
results = {}

# Run all models WITH G1 & G2
print("=" * 60)
print("MODELS WITH G1 & G2 FEATURES")
print("=" * 60)

# Logistic Regression - Base & Tuned
lr_base_acc, lr_base_prec, lr_base_rec, lr_base_f1 = logistic_regression_base(
    X_train_with_grades, X_test_with_grades, "With G1 & G2"
)
print("_" * 50)

lr_tuned_acc, lr_tuned_prec, lr_tuned_rec, lr_tuned_f1 =
↪logistic_regression_tuned(
    X_train_with_grades, X_test_with_grades, "With G1 & G2"
)
print("_" * 50)

# Random Forest - Base & Tuned
rf_base_acc, rf_base_prec, rf_base_rec, rf_base_f1 = random_forest_base(
    X_train_with_grades, X_test_with_grades, "With G1 & G2"
)
print("_" * 50)

```

```

rf_tuned_acc, rf_tuned_prec, rf_tuned_rec, rf_tuned_f1 = random_forest_tuned(
    X_train_with_grades, X_test_with_grades, "With G1 & G2"
)
print("_" * 50)

# Gradient Boosting - Base & Tuned
gbm_base_acc, gbm_base_prec, gbm_base_rec, gbm_base_f1 = gradient_boosting_base(
    X_train_with_grades, X_test_with_grades, "With G1 & G2"
)
print("_" * 50)

gbm_tuned_acc, gbm_tuned_prec, gbm_tuned_rec, gbm_tuned_f1 =
    ↪gradient_boosting_tuned(
        X_train_with_grades, X_test_with_grades, "With G1 & G2"
    )
print("_" * 50)

# XGBoost - Base & Tuned
xgb_base_acc, xgb_base_prec, xgb_base_rec, xgb_base_f1 = xgboost_base(
    X_train_with_grades, X_test_with_grades, "With G1 & G2"
)
print("_" * 50)

xgb_tuned_acc, xgb_tuned_prec, xgb_tuned_rec, xgb_tuned_f1 = xgboost_tuned(
    X_train_with_grades, X_test_with_grades, "With G1 & G2"
)
print("_" * 50)

# SVM - Base & Tuned
svm_base_acc, svm_base_prec, svm_base_rec, svm_base_f1 = svm_base(
    X_train_with_grades, X_test_with_grades, "With G1 & G2"
)
print("_" * 50)

svm_tuned_acc, svm_tuned_prec, svm_tuned_rec, svm_tuned_f1 = svm_tuned(
    X_train_with_grades, X_test_with_grades, "With G1 & G2"
)

# Store results WITH grades
results["with_grades"] = {
    "LR_Base": [lr_base_acc, lr_base_prec, lr_base_rec, lr_base_f1],
    "LR_Tuned": [lr_tuned_acc, lr_tuned_prec, lr_tuned_rec, lr_tuned_f1],
    "RF_Base": [rf_base_acc, rf_base_prec, rf_base_rec, rf_base_f1],
    "RF_Tuned": [rf_tuned_acc, rf_tuned_prec, rf_tuned_rec, rf_tuned_f1],
    "GBM_Base": [gbm_base_acc, gbm_base_prec, gbm_base_rec, gbm_base_f1],
    "GBM_Tuned": [gbm_tuned_acc, gbm_tuned_prec, gbm_tuned_rec, gbm_tuned_f1],
    "XGB_Base": [xgb_base_acc, xgb_base_prec, xgb_base_rec, xgb_base_f1],

```

```

    "XGB_Tuned": [xgb_tuned_acc, xgb_tuned_prec, xgb_tuned_rec, xgb_tuned_f1],
    "SVM_Base": [svm_base_acc, svm_base_prec, svm_base_rec, svm_base_f1],
    "SVM_Tuned": [svm_tuned_acc, svm_tuned_prec, svm_tuned_rec, svm_tuned_f1],
}

```

```

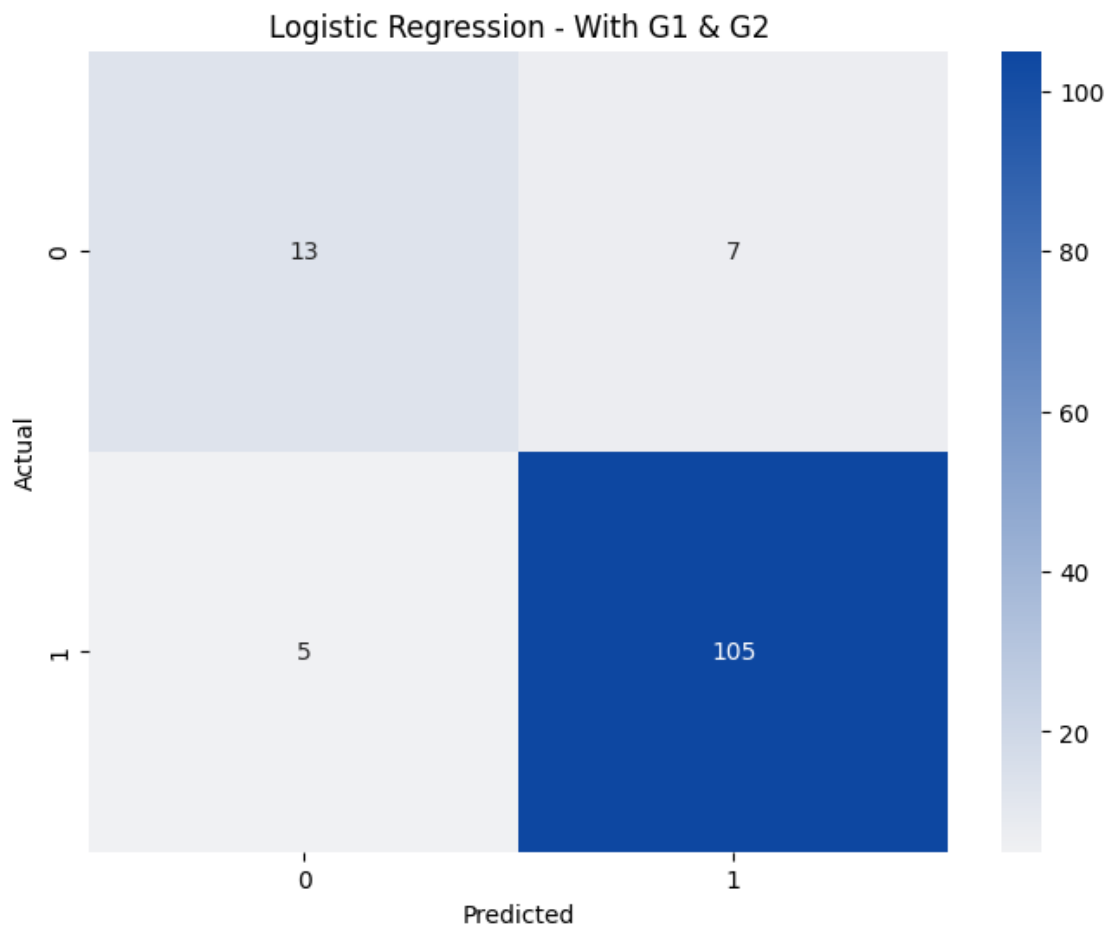
=====
MODELS WITH G1 & G2 FEATURES
=====
LOGISTIC REGRESSION MODEL WITH G1 & G2

```

Classification Report:

	precision	recall	f1-score	support
0	0.72	0.65	0.68	20
1	0.94	0.95	0.95	110
accuracy			0.91	130
macro avg	0.83	0.80	0.82	130
weighted avg	0.90	0.91	0.91	130

Confusion Matrix:



HYPERPARAMETER TUNED LOGISTIC REGRESSION MODEL WITH G1 & G2


Best Hyperparameters for Logistic Regression:

`{'C': 1, 'penalty': 'l1', 'solver': 'liblinear'}`

Classification Report for Improved Logistic Regression:

	precision	recall	f1-score	support
0	0.70	0.70	0.70	20
1	0.95	0.95	0.95	110
accuracy			0.91	130
macro avg	0.82	0.82	0.82	130
weighted avg	0.91	0.91	0.91	130

Confusion Matrix for Improved Model:

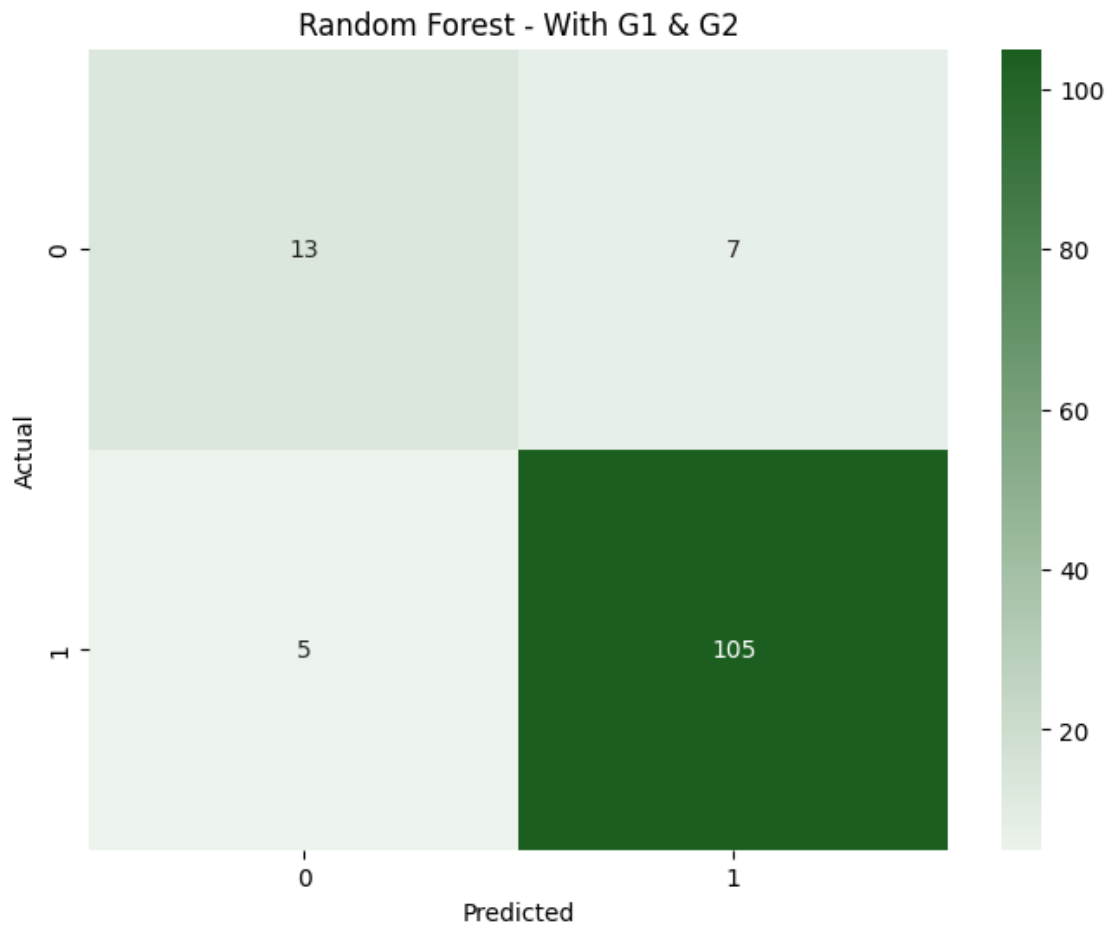


RANDOM FOREST MODEL WITH G1 & G2

Classification Report:

	precision	recall	f1-score	support
0	0.72	0.65	0.68	20
1	0.94	0.95	0.95	110
accuracy			0.91	130
macro avg	0.83	0.80	0.82	130
weighted avg	0.90	0.91	0.91	130

Confusion Matrix:



HYPERPARAMETER TUNED RANDOM FOREST MODEL WITH G1 & G2


Best Hyperparameters for Random Forest:

```
{'n_estimators': 100, 'min_samples_split': 2, 'min_samples_leaf': 4, 'max_depth': 10, 'bootstrap': True}
```

Classification Report for Tuned Random Forest:

	precision	recall	f1-score	support
0	0.69	0.55	0.61	20
1	0.92	0.95	0.94	110
accuracy			0.89	130
macro avg	0.80	0.75	0.77	130
weighted avg	0.89	0.89	0.89	130

Confusion Matrix for Tuned Model:

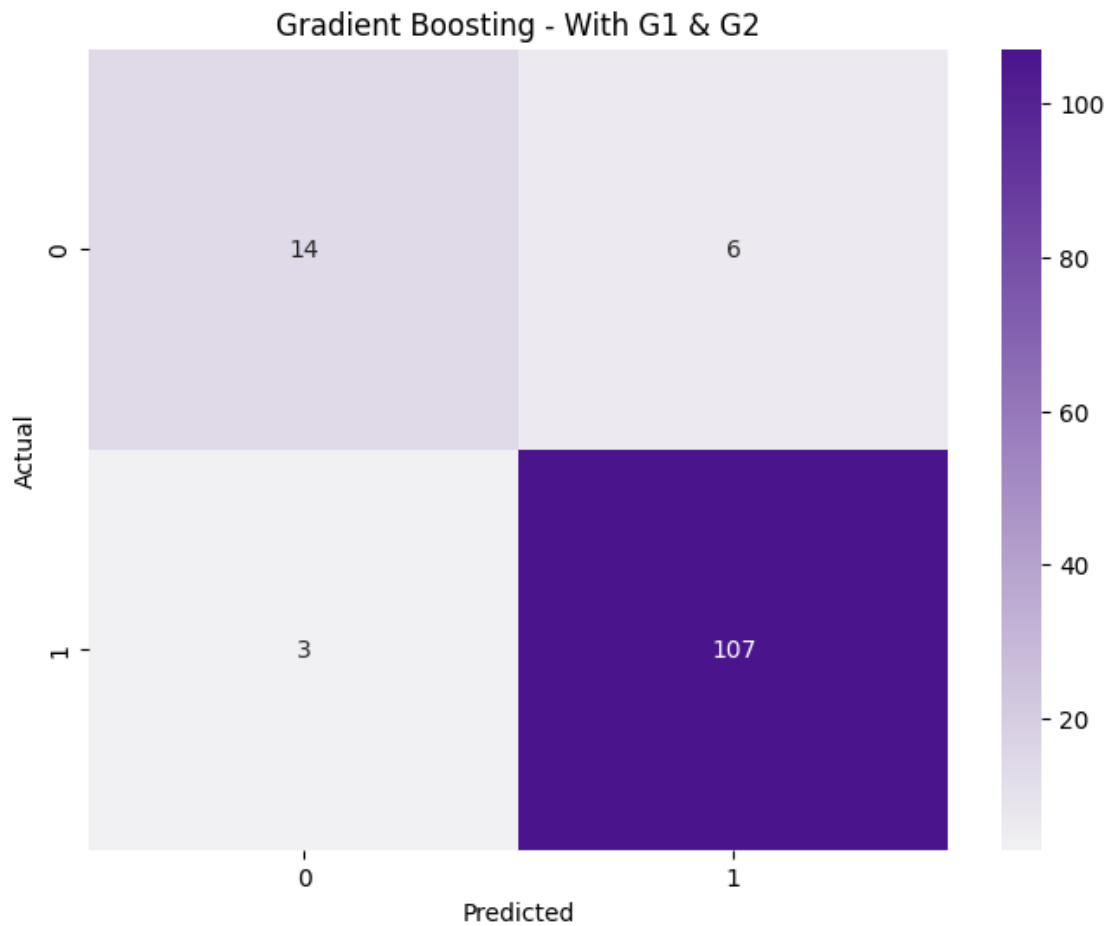


GRADIENT BOOSTING MODEL WITH G1 & G2

Classification Report:

	precision	recall	f1-score	support
0	0.82	0.70	0.76	20
1	0.95	0.97	0.96	110
accuracy			0.93	130
macro avg	0.89	0.84	0.86	130
weighted avg	0.93	0.93	0.93	130

Confusion Matrix:



TUNED GRADIENT BOOSTING MODEL WITH G1 & G2


Best Hyperparameters for Gradient Boosting:

```
{'subsample': 1.0, 'n_estimators': 200, 'min_samples_split': 2,
  'min_samples_leaf': 4, 'max_depth': 7, 'learning_rate': 0.1}
```

Classification Report for Tuned Gradient Boosting:

	precision	recall	f1-score	support
0	0.74	0.70	0.72	20
1	0.95	0.95	0.95	110
accuracy			0.92	130
macro avg	0.84	0.83	0.83	130
weighted avg	0.91	0.92	0.91	130

Confusion Matrix for Tuned GBM:

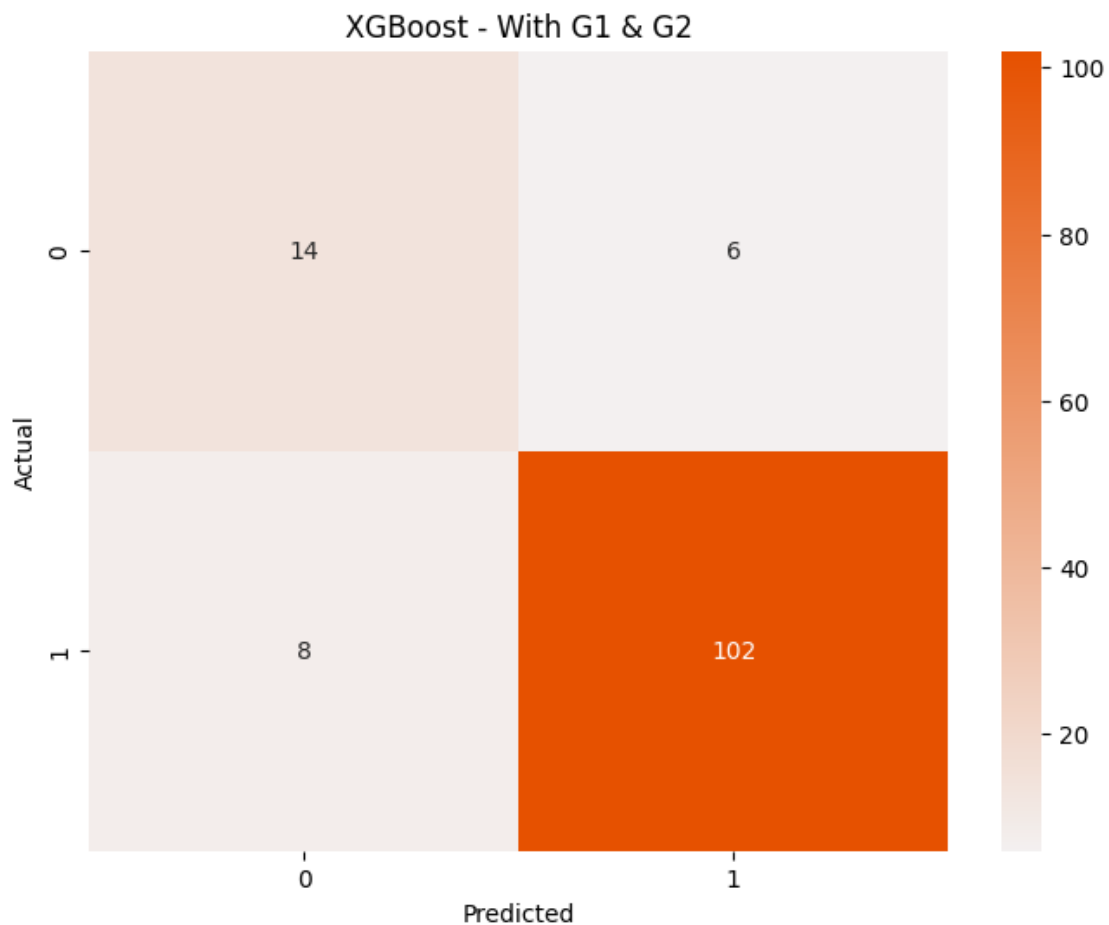


XGBOOST MODEL WITH G1 & G2

Classification Report:

	precision	recall	f1-score	support
0	0.64	0.70	0.67	20
1	0.94	0.93	0.94	110
accuracy			0.89	130
macro avg	0.79	0.81	0.80	130
weighted avg	0.90	0.89	0.89	130

Confusion Matrix:



----- TUNED XGBOOST MODEL WITH G1 & G2

Best Hyperparameters for XGBoost:

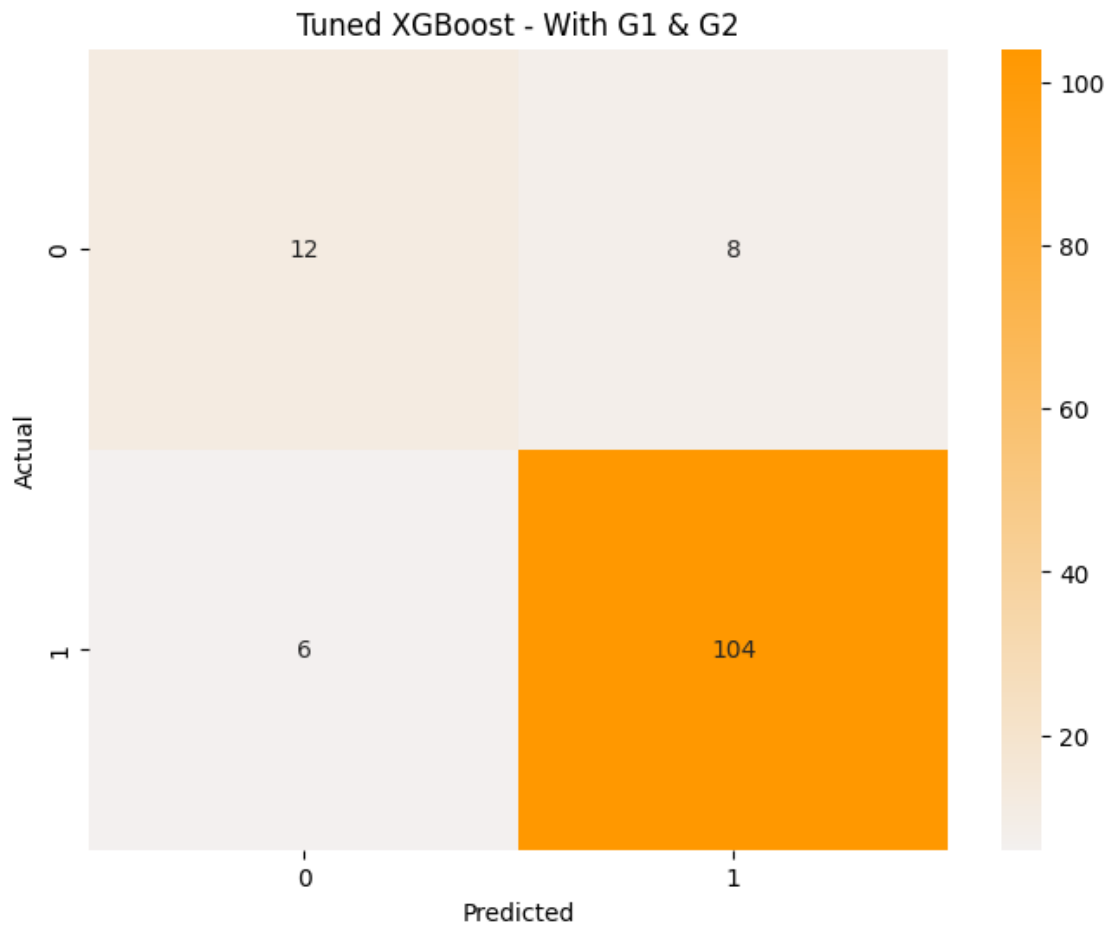
```
{'subsample': 0.8, 'reg_lambda': 1.5, 'reg_alpha': 0, 'n_estimators': 100,  
'min_child_weight': 1, 'max_depth': 3, 'learning_rate': 0.1, 'colsample_bytree':
```

0.9}

Classification Report for Tuned XGBoost:

	precision	recall	f1-score	support
0	0.67	0.60	0.63	20
1	0.93	0.95	0.94	110
accuracy			0.89	130
macro avg	0.80	0.77	0.78	130
weighted avg	0.89	0.89	0.89	130

Confusion Matrix for Tuned XGBoost:




SVM MODEL WITH G1 & G2

Classification Report:

	precision	recall	f1-score	support
0	0.64	0.35	0.45	20
1	0.89	0.96	0.93	110
accuracy			0.87	130
macro avg	0.76	0.66	0.69	130
weighted avg	0.85	0.87	0.85	130

Confusion Matrix:



TUNED SVM MODEL WITH G1 & G2


Best Hyperparameters for SVM:

`{'kernel': 'sigmoid', 'gamma': 0.01, 'degree': 3, 'C': 10}`

Classification Report for Tuned SVM:

	precision	recall	f1-score	support
0	0.67	0.70	0.68	20
1	0.94	0.94	0.94	110
accuracy			0.90	130
macro avg	0.81	0.82	0.81	130
weighted avg	0.90	0.90	0.90	130

Confusion Matrix for Tuned SVM:



```
[12]: # Run all models WITHOUT G1 & G2
print("\n" + "=" * 60)
print("MODELS WITHOUT G1 & G2 FEATURES")
print("=" * 60)

# Logistic Regression - Base & Tuned
```

```

lr_base_acc_no, lr_base_prec_no, lr_base_rec_no, lr_base_f1_no = (
    logistic_regression_base(
        X_train_without_grades, X_test_without_grades, "Without G1 & G2"
    )
)
print("_" * 50)

lr_tuned_acc_no, lr_tuned_prec_no, lr_tuned_rec_no, lr_tuned_f1_no = (
    logistic_regression_tuned(
        X_train_without_grades, X_test_without_grades, "Without G1 & G2"
    )
)
print("_" * 50)

# Random Forest - Base & Tuned
rf_base_acc_no, rf_base_prec_no, rf_base_rec_no, rf_base_f1_no =
    ↪random_forest_base(
        X_train_without_grades, X_test_without_grades, "Without G1 & G2"
    )
print("_" * 50)

rf_tuned_acc_no, rf_tuned_prec_no, rf_tuned_rec_no, rf_tuned_f1_no = (
    random_forest_tuned(
        X_train_without_grades, X_test_without_grades, "Without G1 & G2"
    )
)
print("_" * 50)

# Gradient Boosting - Base & Tuned
gbm_base_acc_no, gbm_base_prec_no, gbm_base_rec_no, gbm_base_f1_no = (
    gradient_boosting_base(
        X_train_without_grades, X_test_without_grades, "Without G1 & G2"
    )
)
print("_" * 50)

gbm_tuned_acc_no, gbm_tuned_prec_no, gbm_tuned_rec_no, gbm_tuned_f1_no = (
    gradient_boosting_tuned(
        X_train_without_grades, X_test_without_grades, "Without G1 & G2"
    )
)
print("_" * 50)

# XGBoost - Base & Tuned
xgb_base_acc_no, xgb_base_prec_no, xgb_base_rec_no, xgb_base_f1_no =
    ↪xgboost_base(
        X_train_without_grades, X_test_without_grades, "Without G1 & G2"
    )

```

```

)
print("_" * 50)

xgb_tuned_acc_no, xgb_tuned_prec_no, xgb_tuned_rec_no, xgb_tuned_f1_no =
    ↪xgb_tuned(
        X_train_without_grades, X_test_without_grades, "Without G1 & G2"
    )
print("_" * 50)

# SVM - Base & Tuned
svm_base_acc_no, svm_base_prec_no, svm_base_rec_no, svm_base_f1_no = svm_base(
    X_train_without_grades, X_test_without_grades, "Without G1 & G2"
)
print("_" * 50)

svm_tuned_acc_no, svm_tuned_prec_no, svm_tuned_rec_no, svm_tuned_f1_no =
    ↪svm_tuned(
        X_train_without_grades, X_test_without_grades, "Without G1 & G2"
    )

# Store results WITHOUT grades
results["without_grades"] = {
    "LR_Base": [lr_base_acc_no, lr_base_prec_no, lr_base_rec_no, lr_base_f1_no],
    "LR_Tuned": [lr_tuned_acc_no, lr_tuned_prec_no, lr_tuned_rec_no,
    ↪lr_tuned_f1_no],
    "RF_Base": [rf_base_acc_no, rf_base_prec_no, rf_base_rec_no, rf_base_f1_no],
    "RF_Tuned": [rf_tuned_acc_no, rf_tuned_prec_no, rf_tuned_rec_no,
    ↪rf_tuned_f1_no],
    "GBM_Base": [gbm_base_acc_no, gbm_base_prec_no, gbm_base_rec_no,
    ↪gbm_base_f1_no],
    "GBM_Tuned": [
        gbm_tuned_acc_no,
        gbm_tuned_prec_no,
        gbm_tuned_rec_no,
        gbm_tuned_f1_no,
    ],
    "XGB_Base": [xgb_base_acc_no, xgb_base_prec_no, xgb_base_rec_no,
    ↪xgb_base_f1_no],
    "XGB_Tuned": [
        xgb_tuned_acc_no,
        xgb_tuned_prec_no,
        xgb_tuned_rec_no,
        xgb_tuned_f1_no,
    ],
    "SVM_Base": [svm_base_acc_no, svm_base_prec_no, svm_base_rec_no,
    ↪svm_base_f1_no],

```

```

    "SVM_Tuned": [
        svm_tuned_acc_no,
        svm_tuned_prec_no,
        svm_tuned_rec_no,
        svm_tuned_f1_no,
    ],
}

```

```

=====
MODELS WITHOUT G1 & G2 FEATURES
=====

```

```

LOGISTIC REGRESSION MODEL WITHOUT G1 & G2

```

```

Classification Report:


```

	precision	recall	f1-score	support
0	0.29	0.25	0.27	20
1	0.87	0.89	0.88	110
accuracy			0.79	130
macro avg	0.58	0.57	0.57	130
weighted avg	0.78	0.79	0.79	130

```

Confusion Matrix:

```



HYPERPARAMETER TUNED LOGISTIC REGRESSION MODEL WITHOUT G1 & G2

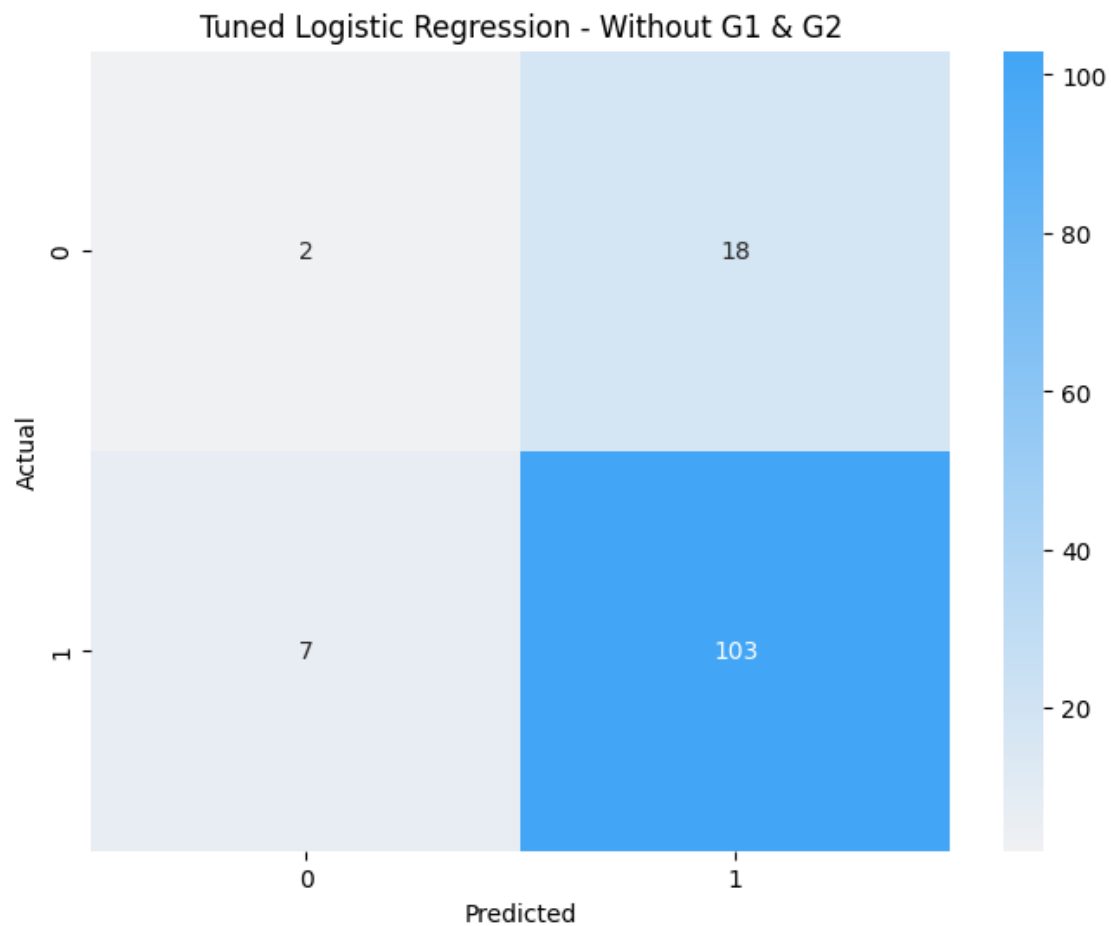
Best Hyperparameters for Logistic Regression:

`{'C': 0.1, 'penalty': 'l1', 'solver': 'saga'}`

Classification Report for Improved Logistic Regression:

	precision	recall	f1-score	support
0	0.22	0.10	0.14	20
1	0.85	0.94	0.89	110
accuracy			0.81	130
macro avg	0.54	0.52	0.51	130
weighted avg	0.75	0.81	0.78	130

Confusion Matrix for Improved Model:

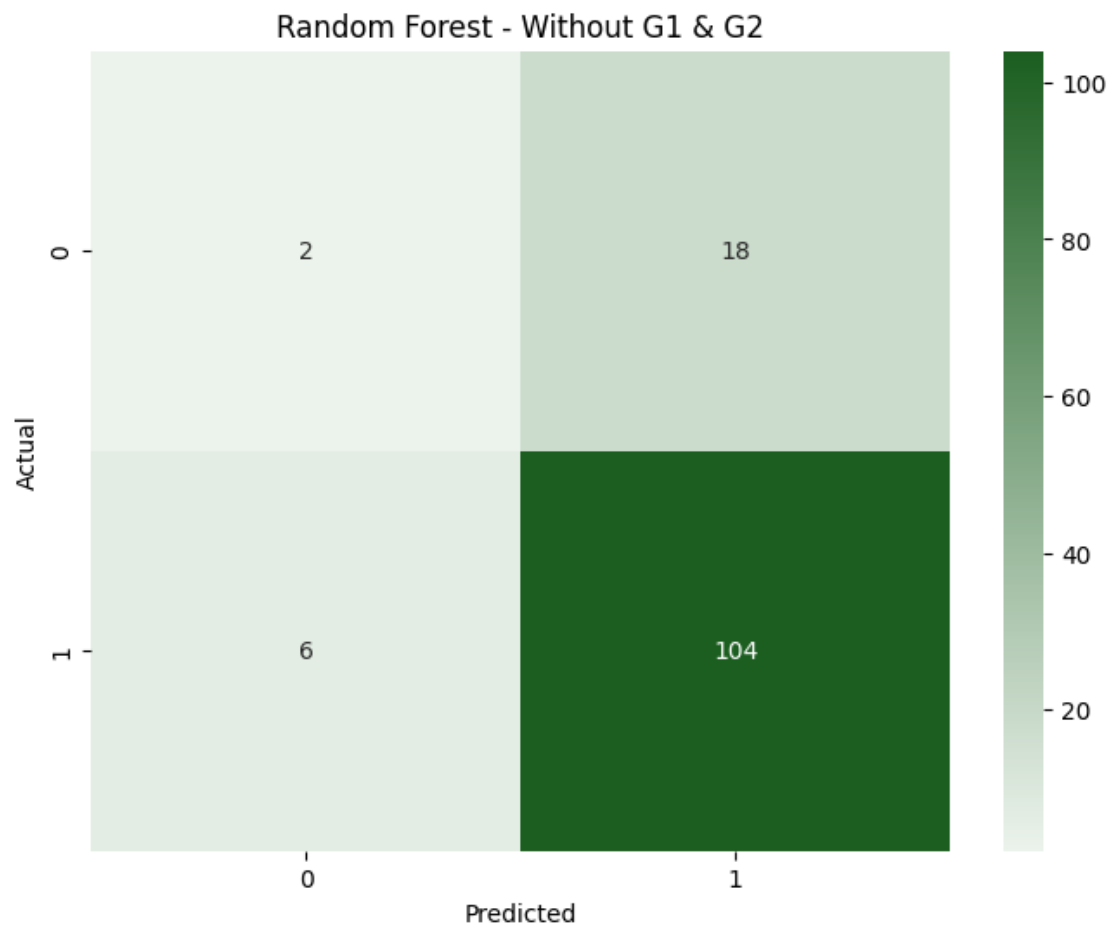


RANDOM FOREST MODEL WITHOUT G1 & G2

Classification Report:

	precision	recall	f1-score	support
0	0.25	0.10	0.14	20
1	0.85	0.95	0.90	110
accuracy			0.82	130
macro avg	0.55	0.52	0.52	130
weighted avg	0.76	0.82	0.78	130

Confusion Matrix:



HYPERPARAMETER TUNED RANDOM FOREST MODEL WITHOUT G1 & G2

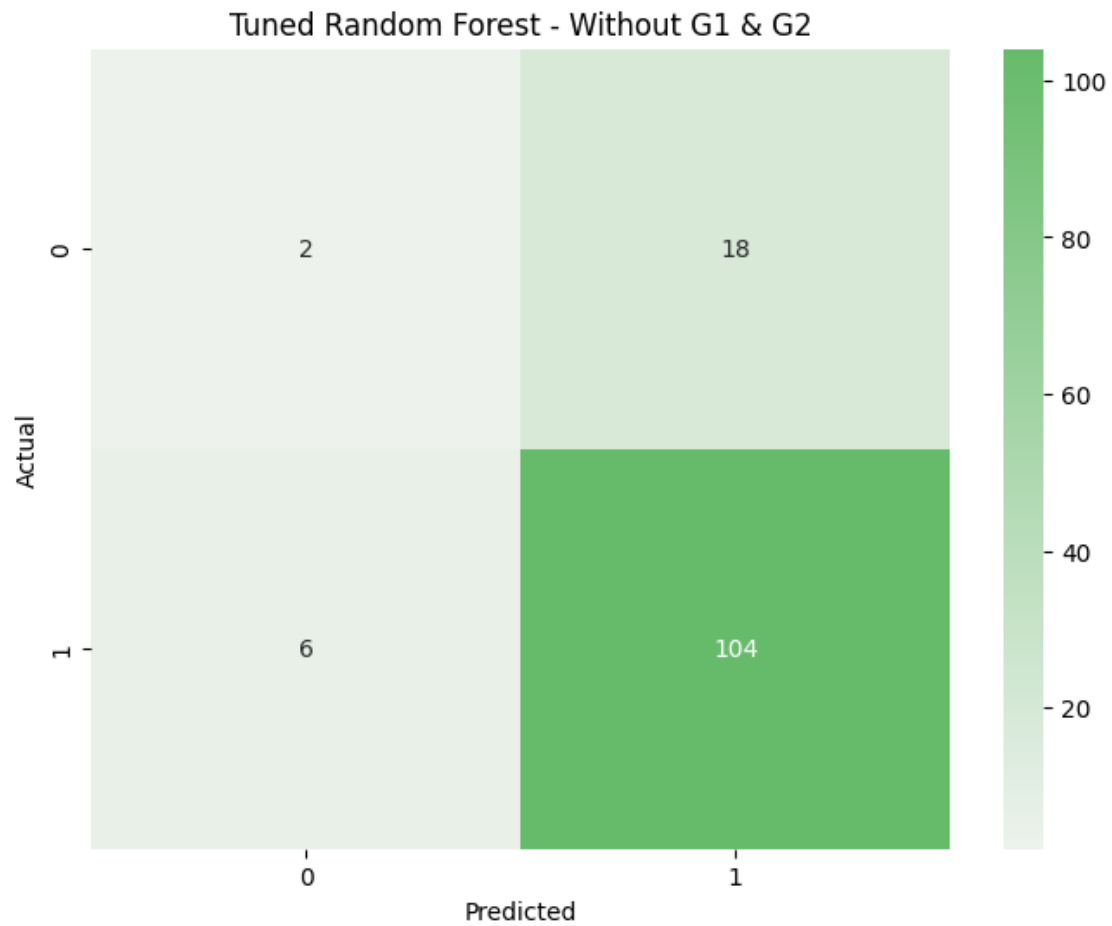
Best Hyperparameters for Random Forest:

```
{'n_estimators': 200, 'min_samples_split': 5, 'min_samples_leaf': 2,
'max_depth': 30, 'bootstrap': True}
```

Classification Report for Tuned Random Forest:

	precision	recall	f1-score	support
0	0.25	0.10	0.14	20
1	0.85	0.95	0.90	110
accuracy			0.82	130
macro avg	0.55	0.52	0.52	130
weighted avg	0.76	0.82	0.78	130

Confusion Matrix for Tuned Model:

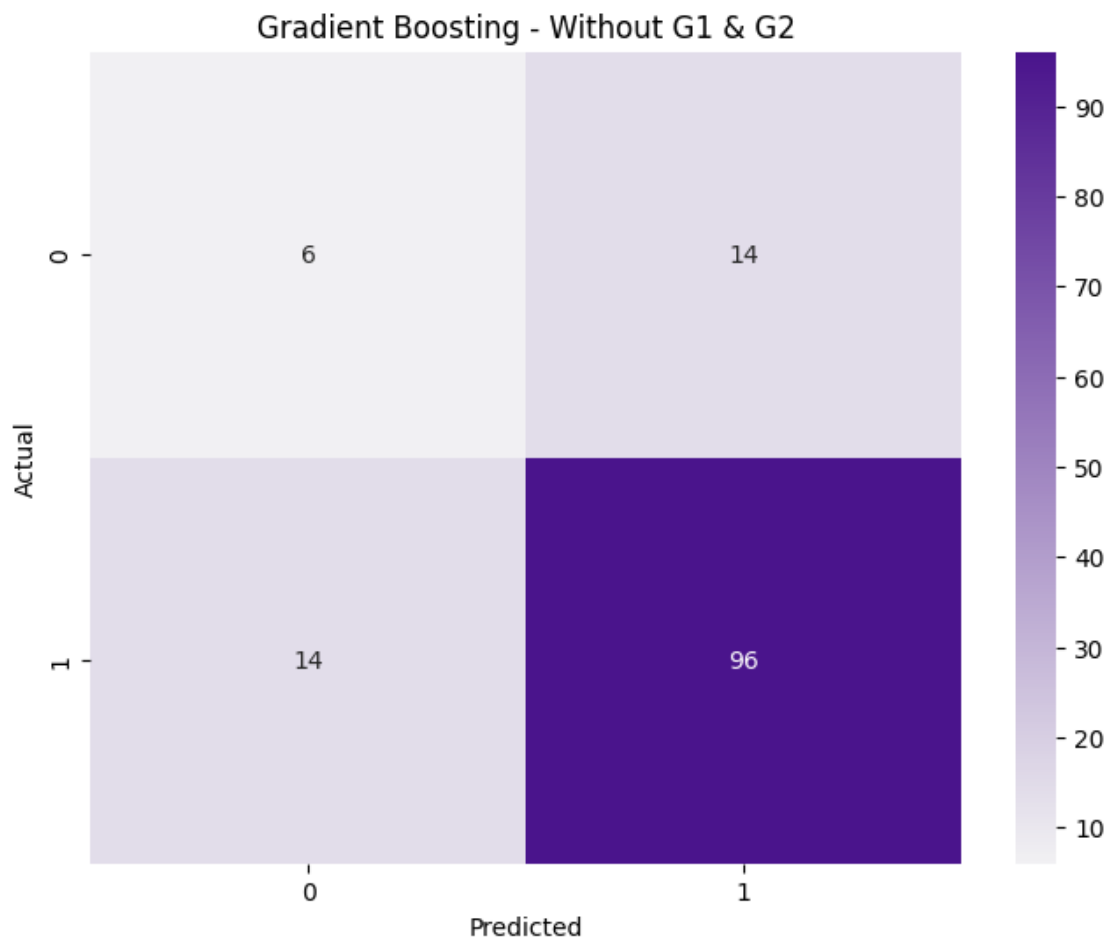


GRADIENT BOOSTING MODEL WITHOUT G1 & G2

Classification Report:

	precision	recall	f1-score	support
0	0.30	0.30	0.30	20
1	0.87	0.87	0.87	110
accuracy			0.78	130
macro avg	0.59	0.59	0.59	130
weighted avg	0.78	0.78	0.78	130

Confusion Matrix:



TUNED GRADIENT BOOSTING MODEL WITHOUT G1 & G2

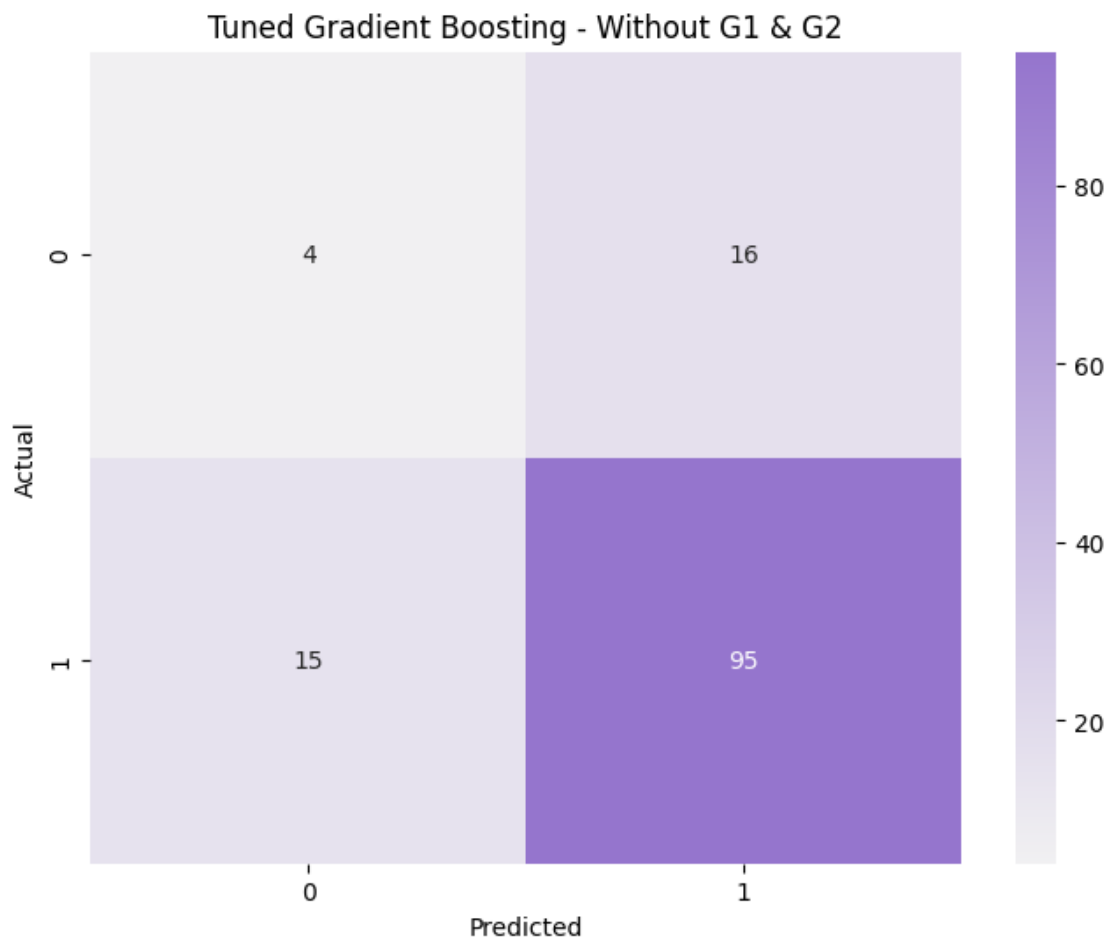
Best Hyperparameters for Gradient Boosting:

```
{'subsample': 1.0, 'n_estimators': 200, 'min_samples_split': 2,
'min_samples_leaf': 4, 'max_depth': 7, 'learning_rate': 0.1}
```

Classification Report for Tuned Gradient Boosting:

	precision	recall	f1-score	support
0	0.21	0.20	0.21	20
1	0.86	0.86	0.86	110
accuracy			0.76	130
macro avg	0.53	0.53	0.53	130
weighted avg	0.76	0.76	0.76	130

Confusion Matrix for Tuned GBM:




XGBOOST MODEL WITHOUT G1 & G2

Classification Report:

	precision	recall	f1-score	support
0	0.25	0.20	0.22	20
1	0.86	0.89	0.88	110
accuracy			0.78	130
macro avg	0.55	0.55	0.55	130
weighted avg	0.77	0.78	0.77	130

Confusion Matrix:



TUNED XGBOOST MODEL WITHOUT G1 & G2


Best Hyperparameters for XGBoost:

```
{'subsample': 1.0, 'reg_lambda': 2, 'reg_alpha': 0.1, 'n_estimators': 200,
'min_child_weight': 1, 'max_depth': 5, 'learning_rate': 0.1, 'colsample_bytree':
1.0}
```

Classification Report for Tuned XGBoost:

	precision	recall	f1-score	support
0	0.25	0.20	0.22	20
1	0.86	0.89	0.88	110
accuracy			0.78	130
macro avg	0.55	0.55	0.55	130
weighted avg	0.77	0.78	0.77	130

Confusion Matrix for Tuned XGBoost:




SVM MODEL WITHOUT G1 & G2

Classification Report:

	precision	recall	f1-score	support
0	0.25	0.10	0.14	20
1	0.85	0.95	0.90	110
accuracy			0.82	130
macro avg	0.55	0.52	0.52	130
weighted avg	0.76	0.82	0.78	130

Confusion Matrix:



TUNED SVM MODEL WITHOUT G1 & G2

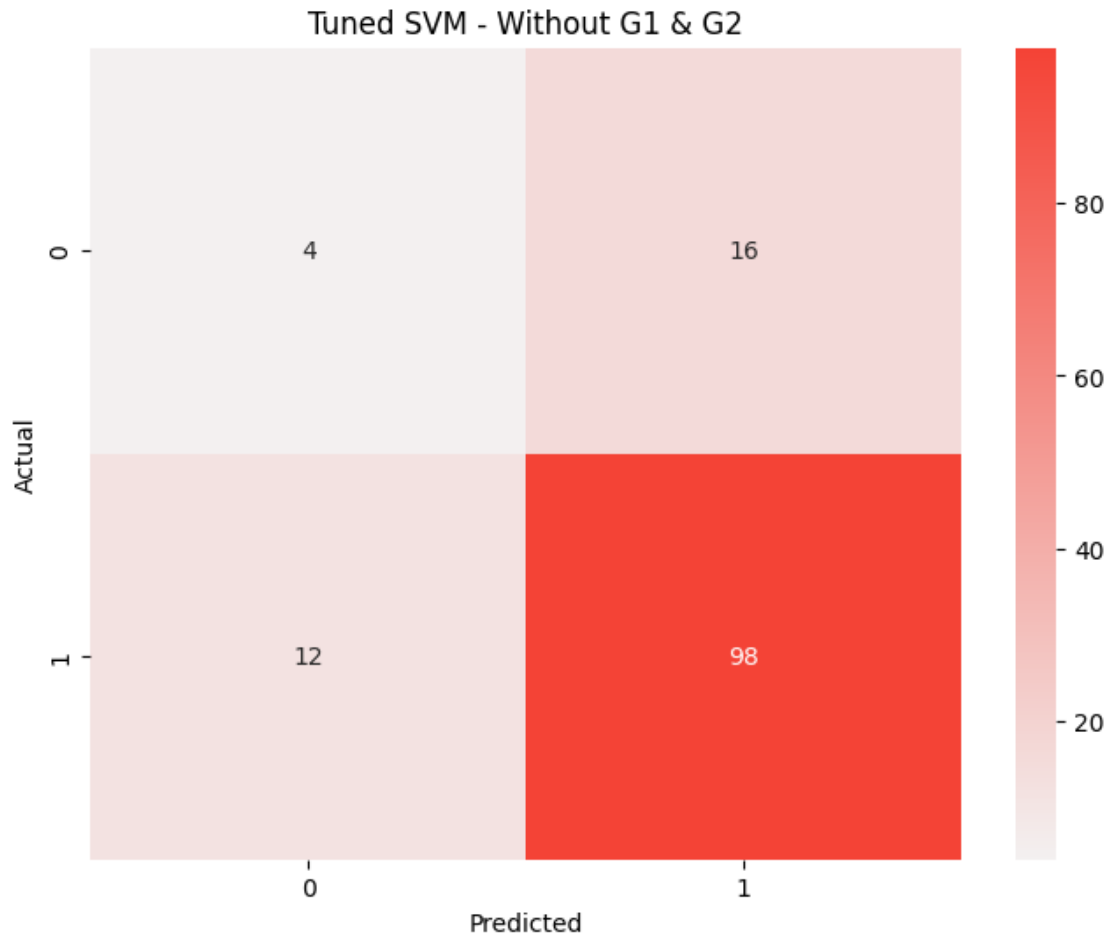
Best Hyperparameters for SVM:

`{'kernel': 'sigmoid', 'gamma': 0.01, 'degree': 3, 'C': 10}`

Classification Report for Tuned SVM:

	precision	recall	f1-score	support
0	0.25	0.20	0.22	20
1	0.86	0.89	0.88	110
accuracy			0.78	130
macro avg	0.55	0.55	0.55	130
weighted avg	0.77	0.78	0.77	130

Confusion Matrix for Tuned SVM:



```
[13]: # Create comprehensive comparison DataFrame
print("\n" + "=" * 80)
print("COMPREHENSIVE MODEL COMPARISON")
print("=" * 80)

# Create comparison dataframes
metrics = ["Accuracy", "Precision", "Recall", "F1-Score"]

# WITH G1 & G2 DataFrame
df_with_grades = pd.DataFrame(results["with_grades"], index=metrics).T
df_with_grades.columns = [f"{col}_WITH_GRADES" for col in df_with_grades.
    ↪columns]

# WITHOUT G1 & G2 DataFrame
df_without_grades = pd.DataFrame(results["without_grades"], index=metrics).T
df_without_grades.columns = [
    f"{col}_WITHOUT_GRADES" for col in df_without_grades.columns]
```

```

]

# Combine both dataframes
comparison_df = pd.concat([df_with_grades, df_without_grades], axis=1)

# Round to 4 decimal places for better readability
comparison_df = comparison_df.round(4)

print("Full Model Comparison Table:")
print(comparison_df)

# Create a focused comparison (F1-Score only)
f1_comparison = pd.DataFrame(
    {
        "With G1 & G2": [
            results["with_grades"][model][3] for model in
            results["with_grades"].keys()
        ],
        "Without G1 & G2": [
            results["without_grades"][model][3]
            for model in results["without_grades"].keys()
        ],
    },
    index=list(results["with_grades"].keys()),
)

print("\nF1-Score Comparison:")
print(f1_comparison.round(4))

# Visualize F1-Score comparison with horizontal bars (sorted)
# Sort by "With G1 & G2" performance (descending)
f1_comparison_sorted = f1_comparison.sort_values(
    "With G1 & G2", ascending=True
) # ascending=True for horizontal bars

plt.figure(figsize=(12, 8))
y_pos = np.arange(len(f1_comparison_sorted.index))
bar_height = 0.35

plt.barh(
    y_pos - bar_height / 2,
    f1_comparison_sorted["With G1 & G2"],
    bar_height,
    label="With G1 & G2",
    alpha=0.8,
    color="#2196f3",
)

```

```

plt.barh(
    y_pos + bar_height / 2,
    f1_comparison_sorted["Without G1 & G2"],
    bar_height,
    label="Without G1 & G2",
    alpha=0.8,
    color="#ff9800",
)

plt.ylabel("Models")
plt.xlabel("F1-Score")
plt.title(
    "Model Performance Comparison: With vs Without G1 & G2 (Sorted by_
↳Performance)"
)

plt.yticks(y_pos, f1_comparison_sorted.index)
plt.legend(loc="lower right")
plt.grid(True, alpha=0.3, axis="x")
plt.tight_layout()
plt.show()

# Find best performing models
print("\nBest Performing Models:")
print(
    f"WITH G1 & G2: {f1_comparison['With G1 & G2'].idxmax()} "
    f"(F1: {f1_comparison['With G1 & G2'].max():.4f})"
)
print(
    f"WITHOUT G1 & G2: {f1_comparison['Without G1 & G2'].idxmax()} "
    f"(F1: {f1_comparison['Without G1 & G2'].max():.4f})"
)

# Calculate improvement from G1 & G2
print("\nImprovement from including G1 & G2:")
improvements = f1_comparison["With G1 & G2"] - f1_comparison["Without G1 & G2"]
for model in improvements.index:
    improvement = improvements[model]
    pct_improvement = (improvement / f1_comparison.loc[model, "Without G1 &
↳G2"]) * 100
    print(f"{model}: +{improvement:.4f} ({pct_improvement:.1f}% improvement)")

```

```

=====
COMPREHENSIVE MODEL COMPARISON
=====

```

Full Model Comparison Table:

	Accuracy_WITH_GRADES	Precision_WITH_GRADES	Recall_WITH_GRADES	\
--	----------------------	-----------------------	--------------------	---

LR_Base	0.9077	0.9375	0.9545
LR_Tuned	0.9077	0.9455	0.9455
RF_Base	0.9077	0.9375	0.9545
RF_Tuned	0.8923	0.9211	0.9545
GBM_Base	0.9308	0.9469	0.9727
GBM_Tuned	0.9154	0.9459	0.9545
XGB_Base	0.8923	0.9444	0.9273
XGB_Tuned	0.8923	0.9286	0.9455
SVM_Base	0.8692	0.8908	0.9636
SVM_Tuned	0.9000	0.9450	0.9364


	F1-Score_WITH_GRADES	Accuracy_WITHOUT_GRADES \
LR_Base	0.9459	0.7923
LR_Tuned	0.9455	0.8077
RF_Base	0.9459	0.8154
RF_Tuned	0.9375	0.8154
GBM_Base	0.9596	0.7846
GBM_Tuned	0.9502	0.7615
XGB_Base	0.9358	0.7846
XGB_Tuned	0.9369	0.7846
SVM_Base	0.9258	0.8154
SVM_Tuned	0.9406	0.7846

	Precision_WITHOUT_GRADES	Recall_WITHOUT_GRADES \
LR_Base	0.8673	0.8909
LR_Tuned	0.8512	0.9364
RF_Base	0.8525	0.9455
RF_Tuned	0.8525	0.9455
GBM_Base	0.8727	0.8727
GBM_Tuned	0.8559	0.8636
XGB_Base	0.8596	0.8909
XGB_Tuned	0.8596	0.8909
SVM_Base	0.8525	0.9455
SVM_Tuned	0.8596	0.8909

	F1-Score_WITHOUT_GRADES
LR_Base	0.8789
LR_Tuned	0.8918
RF_Base	0.8966
RF_Tuned	0.8966
GBM_Base	0.8727
GBM_Tuned	0.8597
XGB_Base	0.8750
XGB_Tuned	0.8750
SVM_Base	0.8966
SVM_Tuned	0.8750

F1-Score Comparison:

	With G1 & G2	Without G1 & G2
LR_Base	0.9459	0.8789
LR_Tuned	0.9455	0.8918
RF_Base	0.9459	0.8966
RF_Tuned	0.9375	0.8966
GBM_Base	0.9596	0.8727
GBM_Tuned	0.9502	0.8597
XGB_Base	0.9358	0.8750
XGB_Tuned	0.9369	0.8750
SVM_Base	0.9258	0.8966
SVM_Tuned	0.9406	0.8750



Best Performing Models:

WITH G1 & G2: GBM_Base (F1: 0.9596)

WITHOUT G1 & G2: RF_Base (F1: 0.8966)

Improvement from including G1 & G2:

LR_Base: +0.0670 (7.6% improvement)

LR_Tuned: +0.0537 (6.0% improvement)

RF_Base: +0.0494 (5.5% improvement)

RF_Tuned: +0.0409 (4.6% improvement)

GBM_Base: +0.0869 (10.0% improvement)

GBM_Tuned: +0.0905 (10.5% improvement)

XGB_Base: +0.0608 (6.9% improvement)
XGB_Tuned: +0.0619 (7.1% improvement)
SVM_Base: +0.0292 (3.3% improvement)
SVM_Tuned: +0.0656 (7.5% improvement)